

CAR PRICE PREDICTION



Here I required to model the price of cars with the available independent variables. It will be used to understand how exactly the prices vary with the independent variables. They can accordingly manipulate the design of the cars, the business strategy etc. to meet certain price levels. Further, the model will be a good way for management to understand the pricing dynamics of a new market.

Here I used 'Car Data' which I got from kaggle.

The independent and dependent variables found in the dataset I used to create this model are displayed in the table below.

DATA DICTONARY		
1	Car_ID	Unique id of each observation (Interger)
2	Symboling	Its assigned insurance risk rating, A value of +3 indicates that the auto is risky, -3 that it is probably pretty safe.(Categorical)
3	carCompany	Name of car company (Categorical)
4	fueltype	Car fuel type i.e gas or diesel (Categorical)
5	aspiration	Aspiration used in a car (Categorical)
6	doornumber	Number of doors in a car (Categorical)
7	carbody	body of car (Categorical)
8	drivewheel	type of drive wheel (Categorical)
9	enginelocation	Location of car engine (Categorical)
10	wheelbase	Weelbase of car (Numeric)
11	carlength	Length of car (Numeric)
12	carwidth	Width of car (Numeric)
13	carheight	height of car (Numeric)
14	curbweight	The weight of a car without occupants or baggage. (Numeric)
15	enginetype	Type of engine. (Categorical)
16	cylindernumber	cylinder placed in the car (Categorical)
17	enginesize	Size of car (Numeric)
18	fuelsystem	Fuel system of car (Categorical)
19	boreratio	Boreratio of car (Numeric)
20	stroke	Stroke or volume inside the engine (Numeric)
21	compressionratio	compression ratio of car (Numeric)
22	horsepower	Horsepower (Numeric)
23	peakrpm	car peak rpm (Numeric)
24	citympg	Mileage in city (Numeric)
25	highwaympg	Mileage on highway (Numeric)
26	price(Dependent variable)	Price of car (Numeric)

Data cleaning and preprocessing

Import Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Read dataset

```
df = pd.read_csv("/content/CarPrice_Assignment.csv")
df
```

	car_ID	symboling	CarName	fueltype	
aspiration \					
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std
..
200	201	-1	volvo 145e (sw)	gas	std
201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo
doornumber	carbody	drivewheel	engine	location	wheelbase ...
\					
0	two	convertible	rwd	front	88.6 ...
1	two	convertible	rwd	front	88.6 ...
2	two	hatchback	rwd	front	94.5 ...
3	four	sedan	fwd	front	99.8 ...
4	four	sedan	4wd	front	99.4 ...

..
200	four	sedan	rwd	front	109.1	...
201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...

	enginesize	fuelsystem	boreratio	stroke	compressionratio
horsepower \					
0	130	mpfi	3.47	2.68	9.0
111					
1	130	mpfi	3.47	2.68	9.0
111					
2	152	mpfi	2.68	3.47	9.0
154					
3	109	mpfi	3.19	3.40	10.0
102					
4	136	mpfi	3.19	3.40	8.0
115					
..
...					
200	141	mpfi	3.78	3.15	9.5
114					
201	141	mpfi	3.78	3.15	8.7
160					
202	173	mpfi	3.58	2.87	8.8
134					
203	145	idi	3.01	3.40	23.0
106					
204	141	mpfi	3.78	3.15	9.5
114					

	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
..
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0

```
204      5400      19      25  22625.0
```

```
[205 rows x 26 columns]
```

Current dataset consist of 205 rows and 26 columns

```
df.head()
```

```
   car_ID  symboling      CarName fueltype aspiration
doornumber \
0         1          3  alfa-romero giulia      gas      std
two
1         2          3  alfa-romero stelvio      gas      std
two
2         3          1  alfa-romero Quadrifoglio      gas      std
two
3         4          2      audi 100 ls      gas      std
four
4         5          2      audi 100ls      gas      std
four
```

```
   carbody drivewheel enginelocation  wheelbase  ...
enginesize \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136
```

```
   fuelsystem  boreratio  stroke  compressionratio  horsepower  peakrpm
citympg \
0      mpfi      3.47      2.68      9.0      111      5000
21
1      mpfi      3.47      2.68      9.0      111      5000
21
2      mpfi      2.68      3.47      9.0      154      5000
19
3      mpfi      3.19      3.40      10.0      102      5500
24
4      mpfi      3.19      3.40      8.0      115      5500
18
```

```
   highwaympg  price
0          27  13495.0
1          27  16500.0
```

```
2      26 16500.0
3      30 13950.0
4      22 17450.0
```

```
[5 rows x 26 columns]
```

```
df.tail()
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber
\						
200	201	-1	volvo 145e (sw)	gas	std	four
201	202	-1	volvo 144ea	gas	turbo	four
202	203	-1	volvo 244dl	gas	std	four
203	204	-1	volvo 246	diesel	turbo	four
204	205	-1	volvo 264gl	gas	turbo	four

	carbody	drivewheel	engine location	wheelbase	...	enginesize
fuelsystem \						
200	sedan	rwd	front	109.1	...	141
mpfi						
201	sedan	rwd	front	109.1	...	141
mpfi						
202	sedan	rwd	front	109.1	...	173
mpfi						
203	sedan	rwd	front	109.1	...	145
idi						
204	sedan	rwd	front	109.1	...	141
mpfi						

	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	\
200	3.78	3.15	9.5	114	5400	23	
201	3.78	3.15	8.7	160	5300	19	
202	3.58	2.87	8.8	134	5500	18	
203	3.01	3.40	23.0	106	4800	26	
204	3.78	3.15	9.5	114	5400	19	

	highwaympg	price
200	28	16845.0
201	25	19045.0
202	23	21485.0
203	27	22470.0
204	25	22625.0

```
[5 rows x 26 columns]
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel           205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase            205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight             205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm               205 non-null    int64
23  citympg               205 non-null    int64
24  highwaympg            205 non-null    int64
25  price                 205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
df.describe()
```

	car_ID	symboling	wheelbase	carlength	carwidth
carheight \					
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805
std	59.322565	1.245307	6.021776	12.337289	2.145204
min	1.000000	-2.000000	86.600000	141.100000	60.300000
25%	52.000000	0.000000	94.500000	166.300000	64.100000
50%	103.000000	1.000000	97.000000	173.200000	65.500000
75%	152.000000	2.000000	101.000000	178.000000	66.000000
max	205.000000	4.000000	112.000000	200.000000	71.000000

```

75%    154.000000    2.000000    102.400000    183.100000    66.900000
55.500000
max     205.000000    3.000000    120.900000    208.100000    72.300000
59.800000

```

```

      curbweight  enginesize  boreratio    stroke
compressionratio \
count    205.000000    205.000000    205.000000    205.000000
205.000000
mean    2555.565854    126.907317     3.329756     3.255415
10.142537
std      520.680204     41.642693     0.270844     0.313597
3.972040
min     1488.000000     61.000000     2.540000     2.070000
7.000000
25%     2145.000000     97.000000     3.150000     3.110000
8.600000
50%     2414.000000    120.000000     3.310000     3.290000
9.000000
75%     2935.000000    141.000000     3.580000     3.410000
9.400000
max     4066.000000    326.000000     3.940000     4.170000
23.000000

```

```

      horsepower      peakrpm      citympg  highwaympg      price
count    205.000000    205.000000    205.000000    205.000000    205.000000
mean     104.117073    5125.121951    25.219512    30.751220    13276.710571
std       39.544167     476.985643     6.542142     6.886443     7988.852332
min       48.000000    4150.000000    13.000000    16.000000     5118.000000
25%       70.000000    4800.000000    19.000000    25.000000     7788.000000
50%       95.000000    5200.000000    24.000000    30.000000    10295.000000
75%      116.000000    5500.000000    30.000000    34.000000    16503.000000
max      288.000000    6600.000000    49.000000    54.000000    45400.000000

```

```
df.dtypes
```

```

car_ID          int64
symboling       int64
CarName         object
fueltype        object
aspiration      object
doornumber      object
carbody         object
drivewheel      object
enginelocation  object
wheelbase       float64
carlength       float64
carwidth        float64
carheight       float64
curbweight      int64

```

```

enginetype      object
cylindernumber  object
enginesize      int64
fuelsystem      object
boreratio       float64
stroke          float64
compressionratio float64
horsepower      int64
peakrpm         int64
citympg         int64
highwaympg      int64
price           float64
dtype: object

```

8 columns are float datatype, 8 columns are integer datatype and 10 columns are object datatype.

```

df.columns
Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
      'doornumber', 'carbody', 'drivewheel', 'enginelocation',
      'wheelbase',
      'carlength', 'carwidth', 'carheight', 'curbweight',
      'enginetype',
      'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio',
      'stroke',
      'compressionratio', 'horsepower', 'peakrpm', 'citympg',
      'highwaympg',
      'price'],
      dtype='object')

```

Dropping duplicate rows

```

df.drop_duplicates()

```

	car_ID	symboling	CarName	fueltype	aspiration	\
0	1	3	alfa-romero giulia	gas	std	
1	2	3	alfa-romero stelvio	gas	std	
2	3	1	alfa-romero Quadrifoglio	gas	std	
3	4	2	audi 100 ls	gas	std	
4	5	2	audi 100ls	gas	std	
..	
200	201	-1	volvo 145e (sw)	gas	std	

201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo

	doornumber	carbody	drivewheel	engine	location	wheelbase	...
\							
0	two	convertible	rwd		front	88.6	...
1	two	convertible	rwd		front	88.6	...
2	two	hatchback	rwd		front	94.5	...
3	four	sedan	fwd		front	99.8	...
4	four	sedan	4wd		front	99.4	...
..
200	four	sedan	rwd		front	109.1	...
201	four	sedan	rwd		front	109.1	...
202	four	sedan	rwd		front	109.1	...
203	four	sedan	rwd		front	109.1	...
204	four	sedan	rwd		front	109.1	...

	enginesize	fuelsystem	boreratio	stroke	compressionratio
horsepower \					
0	130	mpfi	3.47	2.68	9.0
111					
1	130	mpfi	3.47	2.68	9.0
111					
2	152	mpfi	2.68	3.47	9.0
154					
3	109	mpfi	3.19	3.40	10.0
102					
4	136	mpfi	3.19	3.40	8.0
115					
..
...					
200	141	mpfi	3.78	3.15	9.5
114					

201	141	mpfi	3.78	3.15	8.7
160					
202	173	mpfi	3.58	2.87	8.8
134					
203	145	idi	3.01	3.40	23.0
106					
204	141	mpfi	3.78	3.15	9.5
114					

	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
..
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0
204	5400	19	25	22625.0

[205 rows x 26 columns]

Check missing values

```
df.isna().sum()
```

car_ID	0
symboling	0
CarName	0
fueltype	0
aspiration	0
doornumber	0
carbody	0
drivewheel	0
enginelocation	0
wheelbase	0
carlength	0
carwidth	0
carheight	0
curbweight	0
enginetype	0
cylindernumber	0
enginesize	0
fuelsystem	0
boreratio	0
stroke	0
compressionratio	0

```
horsepower      0
peakrpm         0
citympg         0
highwaympg      0
price           0
dtype: int64
```

There are no missing values in the dataset.

Check the 'CarName' feature

```
df['CarName'].unique()
array(['alfa-romero giulia', 'alfa-romero stelvio',
      'alfa-romero Quadrifoglio', 'audi 100 ls', 'audi 100ls',
      'audi fox', 'audi 5000', 'audi 4000', 'audi 5000s (diesel)',
      'bmw 320i', 'bmw x1', 'bmw x3', 'bmw z4', 'bmw x4', 'bmw x5',
      'chevrolet impala', 'chevrolet monte carlo', 'chevrolet vega
2300',
      'dodge rampage', 'dodge challenger se', 'dodge d200',
      'dodge monaco (sw)', 'dodge colt hardtop', 'dodge colt (sw)',
      'dodge coronet custom', 'dodge dart custom',
      'dodge coronet custom (sw)', 'honda civic', 'honda civic cvcc',
      'honda accord cvcc', 'honda accord lx', 'honda civic 1500 gl',
      'honda accord', 'honda civic 1300', 'honda prelude',
      'honda civic (auto)', 'isuzu MU-X', 'isuzu D-Max ',
      'isuzu D-Max V-Cross', 'jaguar xj', 'jaguar xf', 'jaguar xk',
      'maxda rx3', 'maxda glc deluxe', 'mazda rx2 coupe', 'mazda rx-
4',
      'mazda glc deluxe', 'mazda 626', 'mazda glc', 'mazda rx-7 gs',
      'mazda glc 4', 'mazda glc custom l', 'mazda glc custom',
      'buick electra 225 custom', 'buick century luxus (sw)',
      'buick century', 'buick skyhawk', 'buick opel isuzu deluxe',
      'buick skylark', 'buick century special',
      'buick regal sport coupe (turbo)', 'mercury cougar',
      'mitsubishi mirage', 'mitsubishi lancer', 'mitsubishi
outlander',
      'mitsubishi g4', 'mitsubishi mirage g4', 'mitsubishi montero',
      'mitsubishi pajero', 'Nissan versa', 'nissan gt-r', 'nissan
rogue',
      'nissan latio', 'nissan titan', 'nissan leaf', 'nissan juke',
      'nissan note', 'nissan clipper', 'nissan nv200', 'nissan dayz',
      'nissan fuga', 'nissan otti', 'nissan teana', 'nissan kicks',
      'peugeot 504', 'peugeot 304', 'peugeot 504 (sw)', 'peugeot
604sl',
      'peugeot 505s turbo diesel', 'plymouth fury iii',
      'plymouth cricket', 'plymouth satellite custom (sw)',
      'plymouth fury gran sedan', 'plymouth valiant', 'plymouth
duster',
```

```
'porsche macan', 'porsche panamera', 'porsche cayenne',  
'porsche boxster', 'renault 12tl', 'renault 5 gtl', 'saab 99e',  
'saab 99le', 'saab 99gle', 'subaru', 'subaru dl', 'subaru brz',  
'subaru baja', 'subaru r1', 'subaru r2', 'subaru trezia',  
'subaru tribeca', 'toyota corona mark ii', 'toyota corona',  
'toyota corolla 1200', 'toyota corona hardtop',  
'toyota corolla 1600 (sw)', 'toyota carina', 'toyota mark ii',  
'toyota corolla', 'toyota corolla liftback',  
'toyota celica gt liftback', 'toyota corolla tercel',  
'toyota corona liftback', 'toyota starlet', 'toyota tercel',  
'toyota cressida', 'toyota celica gt', 'toyota tercel',  
'volkswagen rabbit', 'volkswagen 1131 deluxe sedan',  
'volkswagen model 111', 'volkswagen type 3', 'volkswagen 411  
(sw)',  
'volkswagen super beetle', 'volkswagen dasher', 'vw dasher',  
'vw rabbit', 'volkswagen rabbit', 'volkswagen rabbit custom',  
'volvo 145e (sw)', 'volvo 144ea', 'volvo 244dl', 'volvo 245',  
'volvo 264gl', 'volvo diesel', 'volvo 246'], dtype=object)
```

Check 'fueltype' feature

```
df['fueltype'].unique()  
array(['gas', 'diesel'], dtype=object)
```

Check 'aspiration' feature

```
df['aspiration'].unique()  
array(['std', 'turbo'], dtype=object)
```

Check 'doornumber' feature

```
df['doornumber'].unique()  
array(['two', 'four'], dtype=object)
```

Check 'carbody' feature

```
df['carbody'].unique()  
array(['convertible', 'hatchback', 'sedan', 'wagon', 'hardtop'],  
      dtype=object)
```

Check 'drivewheel' feature

```
df['drivewheel'].unique()  
array(['rwd', 'fwd', '4wd'], dtype=object)
```

Check 'enginelocation' feature

```
df['enginelocation'].unique()  
array(['front', 'rear'], dtype=object)
```

Check 'enginetype' feature

```
df['enginetype'].unique()  
array(['dohc', 'ohcv', 'ohc', 'l', 'rotor', 'ohcf', 'dohcv'],  
      dtype=object)
```

Check 'cylindernumber' feature

```
df['cylindernumber'].unique()  
array(['four', 'six', 'five', 'three', 'twelve', 'two', 'eight'],  
      dtype=object)
```


Check 'fuelsystem' feature

```
df['fuelsystem'].unique()  
array(['mpfi', '2bbl', 'mfi', '1bbl', 'spfi', '4bbl', 'idi', 'spdi'],  
      dtype=object)
```

Data visualizations

Check the distribution of car price

```
plt.figure(figsize=(16, 10))  
sns.histplot(data=df, x='price', kde=True)  
<Axes: xlabel='price', ylabel='Count'>
```



```
df['price'].describe()
```

```
count      205.000000
mean       13276.710571
std        7988.852332
min         5118.000000
25%        7788.000000
50%       10295.000000
75%       16503.000000
max       45400.000000
Name: price, dtype: float64
```

The plot seemed to be right-skewed, meaning that the most prices in the dataset are low (Below 15,000).

There is a significant difference between the mean and the median of the price distribution.

Minimum price is \$5118

Maximum price is \$5400

Data correlation using Heatmap


Multivariate analysis

```
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(),annot=True,cmap='coolwarm')
plt.title('Heatmap of the dataset')
```

<ipython-input-154-5daefb561f69>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(),annot=True,cmap='coolwarm')
```

```
Text(0.5, 1.0, 'Heatmap of the dataset')
```



Highly correlated variables to price are - carwidth, curbweight, enginesize and horsepower.

Visualising Categorical Data¶

- CompanyName
- fueltype
- enginetype
- carbody
- doornumber

- enginelocation
- fuelsystem
- cylindernumber
- aspiration
- drivewheel

Visualising categorical data

```
plt.figure(figsize=(30,30))

plt.subplot(6,2,1)
sns.countplot(x='CarName',data=df)
plt.title('Count of CarName')

plt.subplot(6,2,2)
sns.countplot(x='symboling',data=df)
plt.title('Count of symboling')

plt.subplot(6,2,3)
sns.countplot(x='fueltype',data=df)
plt.title('Count of fueltype')

plt.subplot(6,2,4)
sns.countplot(x='enginetype',data=df)
plt.title('Count of enginetype')

plt.subplot(6,2,5)
sns.countplot(x='carbody',data=df)
plt.title('Count of carbody')

plt.subplot(6,2,6)
sns.countplot(x='doornumber',data=df)
plt.title('Count of doornumber')

plt.subplot(6,2,7)
sns.countplot(x='enginelocation',data=df)
plt.title('Count of enginelocation')

plt.subplot(6,2,8)
sns.countplot(x='fuelsystem',data=df)
plt.title('Count of fuelsystem')

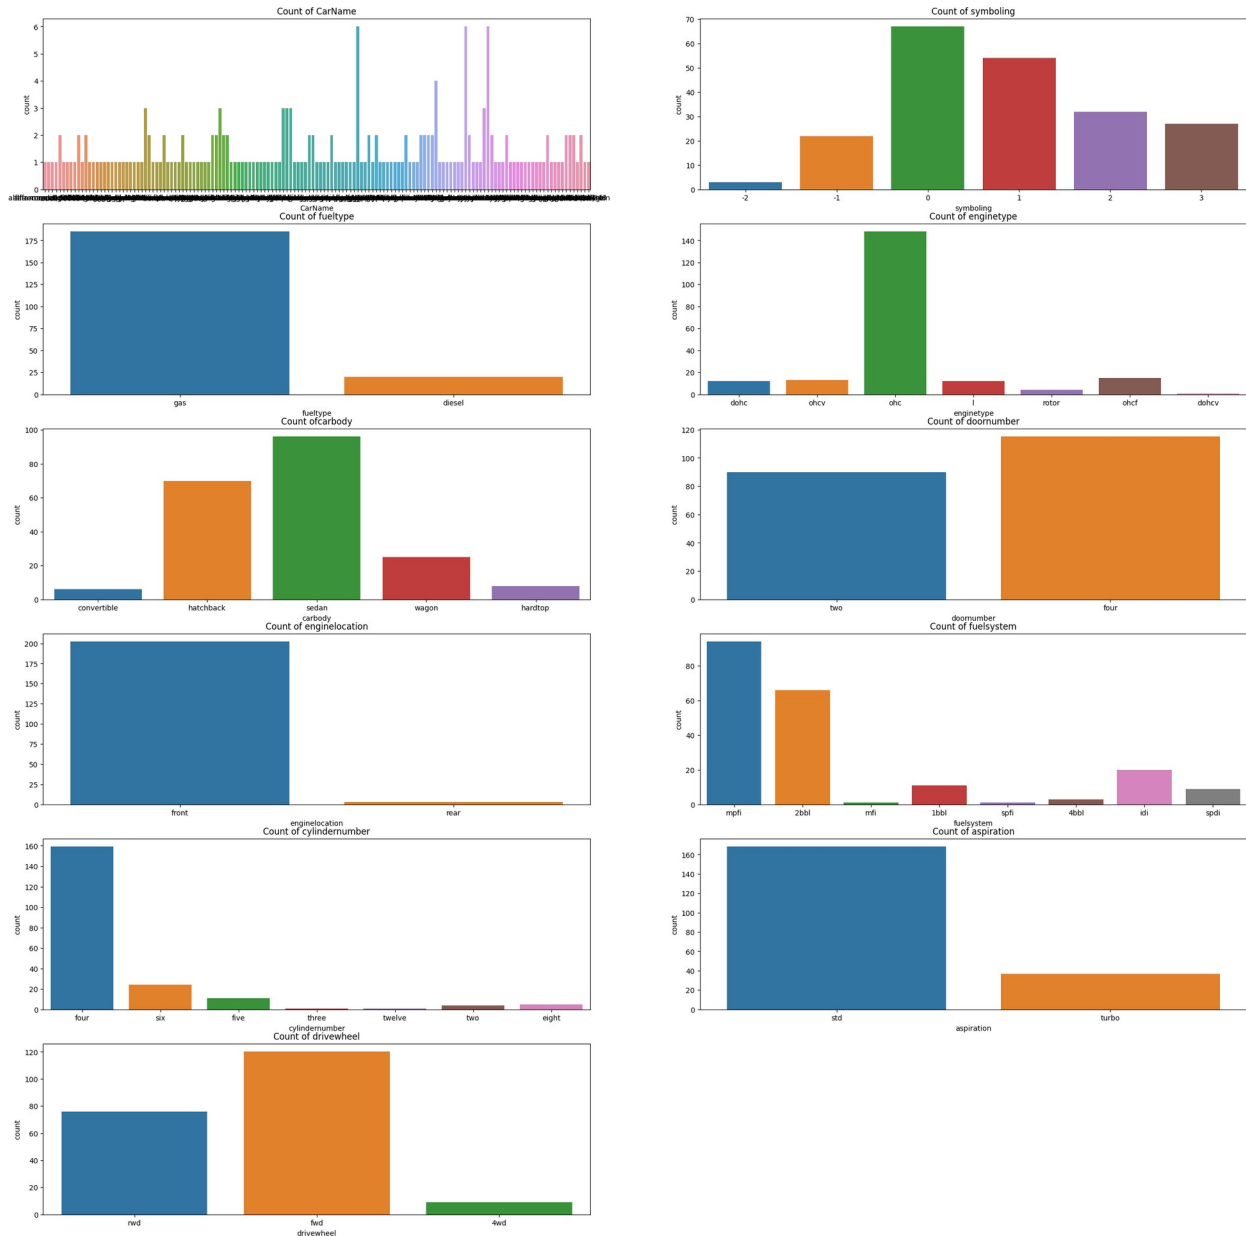
plt.subplot(6,2,9)
sns.countplot(x='cylindernumber',data=df)
plt.title('Count of cylindernumber')

plt.subplot(6,2,10)
sns.countplot(x='aspiration',data=df)
plt.title('Count of aspiration')
```



```
plt.subplot(6,2,11)
sns.countplot(x='drivewheel',data=df)
plt.title('Count of drivewheel')

Text(0.5, 1.0, 'Count of drivewheel')
```



Number of gas fueled cars are more than diesel.

sedan is the top car type preferred.

It seems that the symboling with 0 and 1 values have high number of rows (i.e. They are most sold.)

ohc Engine type seems to be most favored type.

We can see most of the cars use gas as fuel.

FWD is the most favored, and the second place is RWD followed by 4WD.

Most of the cars' engine location are front.

Most of the cars are four-cylinder.


MPFI is the favored type of fuel system.

Visualising numerical variable


symboling
wheelbase
carlength
carwidth
carheight
curbweight
enginesize
boreratio
stroke
compressionratio
horsepower
peakrpm
citympg
highwaympg

```
x=['symboling','wheelbase','carlength','carwidth','carheight','curbweight','enginesize','boreratio','stroke','compressionratio','horsepower','peakrpm','citympg','highwaympg']
for variable in x:
    plt.scatter(df[variable], df['price'])
    plt.xlabel(variable)
    plt.ylabel('Price')
    plt.title(f'Scatter Plot of {variable} vs Price')
    plt.show()
```


Scatter Plot of symboling vs Price




Scatter Plot of wheelbase vs Price




Scatter Plot of carlength vs Price




Scatter Plot of carwidth vs Price




Scatter Plot of carheight vs Price




Scatter Plot of curbweight vs Price




Scatter Plot of enginesize vs Price




Scatter Plot of boreratio vs Price




Scatter Plot of stroke vs Price




Scatter Plot of compressionratio vs Price




Scatter Plot of horsepower vs Price




Scatter Plot of peakrpm vs Price



Scatter Plot of citympg vs Price





carwidth, carlength and curbweight seems to have a poitive correlation with price.

carheight doesn't show any significant trend with price.

enginesize, boreratio, horsepower, wheelbase - seem to have a significant positive correlation with price.

citympg, highwaympg - seem to have a significant negative correlation with price.

Dropping unwanted column

```
df.drop(['car_ID'],axis=1,inplace=True)
df
```

symboling	CarName	fueltype	aspiration
doornumber \ 0	3	alfa-romero giulia	gas std
two			
1	3	alfa-romero stelvio	gas std
two			
2	1	alfa-romero Quadrifoglio	gas std
two			
3	2	audi 100 ls	gas std
four			

4	2	audi 100ls	gas	std		
four						
..
.						
200	-1	volvo 145e (sw)	gas	std		
four						
201	-1	volvo 144ea	gas	turbo		
four						
202	-1	volvo 244dl	gas	std		
four						
203	-1	volvo 246	diesel	turbo		
four						
204	-1	volvo 264gl	gas	turbo		
four						
carbody drivewheel enginelocation wheelbase carlength ...						
\						
0	convertible	rwd	front	88.6	168.8	...
1	convertible	rwd	front	88.6	168.8	...
2	hatchback	rwd	front	94.5	171.2	...
3	sedan	fwd	front	99.8	176.6	...
4	sedan	4wd	front	99.4	176.6	...
..
200	sedan	rwd	front	109.1	188.8	...
201	sedan	rwd	front	109.1	188.8	...
202	sedan	rwd	front	109.1	188.8	...
203	sedan	rwd	front	109.1	188.8	...
204	sedan	rwd	front	109.1	188.8	...
enginesize fuelsystem boreratio stroke compressionratio						
horsepower \						
0	130	mpfi	3.47	2.68		9.0
111						
1	130	mpfi	3.47	2.68		9.0
111						
2	152	mpfi	2.68	3.47		9.0
154						
3	109	mpfi	3.19	3.40		10.0
102						

4	136	mpfi	3.19	3.40	8.0
115					
...
...					
200	141	mpfi	3.78	3.15	9.5
114					
201	141	mpfi	3.78	3.15	8.7
160					
202	173	mpfi	3.58	2.87	8.8
134					
203	145	idi	3.01	3.40	23.0
106					
204	141	mpfi	3.78	3.15	9.5
114					

	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
...
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0
204	5400	19	25	22625.0

[205 rows x 25 columns]

Encoding using LabelEncoder

```
from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
df['CarName']=encoder.fit_transform(df['CarName'])
df['fueltype']=encoder.fit_transform(df['fueltype'])
df['aspiration']=encoder.fit_transform(df['aspiration'])
df['doornumber']=encoder.fit_transform(df['doornumber'])
df['carbody']=encoder.fit_transform(df['carbody'])
df['drivewheel']=encoder.fit_transform(df['drivewheel'])
df['enginelocation']=encoder.fit_transform(df['enginelocation'])
df['enginetype']=encoder.fit_transform(df['enginetype'])
df['cylindernumber']=encoder.fit_transform(df['cylindernumber'])
df['fuelsystem']=encoder.fit_transform(df['fuelsystem'])

df.dtypes
```

symboling	int64
CarName	int64
fueltype	int64

```

aspiration          int64
doornumber          int64
carbody             int64
drivewheel          int64
engineloation       int64
wheelbase           float64
carlength           float64
carwidth            float64
carheight           float64
curbweight          int64
enginetype          int64
cylindernumber      int64
enginesize          int64
fuelsystem          int64
boreratio           float64
stroke              float64
compressionratio    float64
horsepower          int64
peakrpm             int64
citympg             int64
highwaympg          int64
price               float64
dtype: object

```

df

	symboling	CarName	fueltype	aspiration	doornumber	carbody	\
0	3	2	1	0	1	0	
1	3	3	1	0	1	0	
2	1	1	1	0	1	2	
3	2	4	1	0	0	3	
4	2	5	1	0	0	3	
..	
200	-1	139	1	0	0	3	
201	-1	138	1	1	0	3	
202	-1	140	1	0	0	3	
203	-1	142	0	1	0	3	
204	-1	143	1	1	0	3	

	drivewheel	engineloation	wheelbase	carlength	...	enginesize
\						
0	2	0	88.6	168.8	...	130
1	2	0	88.6	168.8	...	130
2	2	0	94.5	171.2	...	152
3	1	0	99.8	176.6	...	109
4	0	0	99.4	176.6	...	136

...
200	2	0	109.1	188.8	...	141
201	2	0	109.1	188.8	...	141
202	2	0	109.1	188.8	...	173
203	2	0	109.1	188.8	...	145
204	2	0	109.1	188.8	...	141
	fuelsystem	boreratio	stroke	compressionratio	horsepower	
peakrpm \						
0	5	3.47	2.68	9.0	111	
5000						
1	5	3.47	2.68	9.0	111	
5000						
2	5	2.68	3.47	9.0	154	
5000						
3	5	3.19	3.40	10.0	102	
5500						
4	5	3.19	3.40	8.0	115	
5500						
...
...						
200	5	3.78	3.15	9.5	114	
5400						
201	5	3.78	3.15	8.7	160	
5300						
202	5	3.58	2.87	8.8	134	
5500						
203	3	3.01	3.40	23.0	106	
4800						
204	5	3.78	3.15	9.5	114	
5400						
	citympg	highwaympg	price			
0	21	27	13495.0			
1	21	27	16500.0			
2	19	26	16500.0			
3	24	30	13950.0			
4	18	22	17450.0			
...			
200	23	28	16845.0			
201	19	25	19045.0			
202	18	23	21485.0			
203	26	27	22470.0			

204 19 25 22625.0

[205 rows x 25 columns]

Separate the data into input and output data

```
x=df.iloc[:, :-1].values
```

```
y=df.iloc[:, -1].values
```

```
y
```

```
array([[13495.   , 16500.   , 16500.   , 13950.   , 17450.   , 15250.   ,
        17710.   , 18920.   , 23875.   , 17859.167, 16430.   , 16925.   ,
        20970.   , 21105.   , 24565.   , 30760.   , 41315.   , 36880.   ,
        5151.   , 6295.   , 6575.   , 5572.   , 6377.   , 7957.   ,
        6229.   , 6692.   , 7609.   , 8558.   , 8921.   , 12964.   ,
        6479.   , 6855.   , 5399.   , 6529.   , 7129.   , 7295.   ,
        7295.   , 7895.   , 9095.   , 8845.   , 10295.   , 12945.   ,
        10345.   , 6785.   , 8916.5 , 8916.5 , 11048.   , 32250.   ,
        35550.   , 36000.   , 5195.   , 6095.   , 6795.   , 6695.   ,
        7395.   , 10945.   , 11845.   , 13645.   , 15645.   , 8845.   ,
        8495.   , 10595.   , 10245.   , 10795.   , 11245.   , 18280.   ,
        18344.   , 25552.   , 28248.   , 28176.   , 31600.   , 34184.   ,
        35056.   , 40960.   , 45400.   , 16503.   , 5389.   , 6189.   ,
        6669.   , 7689.   , 9959.   , 8499.   , 12629.   , 14869.   ,
        14489.   , 6989.   , 8189.   , 9279.   , 9279.   , 5499.   ,
        7099.   , 6649.   , 6849.   , 7349.   , 7299.   , 7799.   ,
        7499.   , 7999.   , 8249.   , 8949.   , 9549.   , 13499.   ,
        14399.   , 13499.   , 17199.   , 19699.   , 18399.   , 11900.   ,
        13200.   , 12440.   , 13860.   , 15580.   , 16900.   , 16695.   ,
        17075.   , 16630.   , 17950.   , 18150.   , 5572.   , 7957.   ])
```

```

',
    6229.    , 6692.    , 7609.    , 8921.    , 12764.    , 22018.
',
    32528.   , 34028.   , 37028.   , 31400.5  , 9295.    , 9895.
',
    11850.   , 12170.   , 15040.   , 15510.   , 18150.   , 18620.
',
    5118.    , 7053.    , 7603.    , 7126.    , 7775.    , 9960.
',
    9233.    , 11259.   , 7463.    , 10198.   , 8013.    , 11694.
',
    5348.    , 6338.    , 6488.    , 6918.    , 7898.    , 8778.
',
    6938.    , 7198.    , 7898.    , 7788.    , 7738.    , 8358.
',
    9258.    , 8058.    , 8238.    , 9298.    , 9538.    , 8449.
',
    9639.    , 9989.    , 11199.   , 11549.   , 17669.   , 8948.
',
    10698.   , 9988.    , 10898.   , 11248.   , 16558.   , 15998.
',
    15690.   , 15750.   , 7775.    , 7975.    , 7995.    , 8195.
',
    8495.    , 9495.    , 9995.    , 11595.   , 9980.    , 13295.
',
    13845.   , 12290.   , 12940.   , 13415.   , 15985.   , 16515.
',
    18420.   , 18950.   , 16845.   , 19045.   , 21485.   , 22470.
',
    22625.   ] )

```

Split the data into training and testing data

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)

```

Scaling using StandardScaler

```

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(x_train)
x_train=scaler.transform(x_train)
x_test=scaler.transform(x_test)

```

Model Creation

```

from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
y_pred

array([25263.71155727, 17510.60100873, 9440.29961597, 13532.18496334,
       26109.69474918, 6312.04809043, 7584.6572441 , 5893.76273471,
       9242.62222642, 5821.86589929, 13857.61067278, 6162.24167276,
       16482.82264907, 10706.77526758, 40393.49162707, 6861.31339325,
       -151.21120815, 14789.44762836, 9670.36006267, 10331.20595444,
       11095.07488034, 20819.20389829, 8195.27918247, 3546.04281395,
       7755.9093543 , 23953.59319483, 14433.67089985, 15762.62290799,
       5067.19507839, 16242.84673396, 26307.25262831, 7381.45646566,
       4235.16070955, 22127.7853912 , 8452.86521883, 26796.96656169,
       10151.62813053, 9736.48156002, 6845.13061313, 15091.05428725,
       7338.31882785, 13461.96663554, 18879.25403184, 4724.11656838,
       6735.85164657, 10078.5772291 , 8866.25914506, 6925.7671163 ,
       18238.42401821, 15330.17170675, 7106.81136173, 19425.30940642,
       3305.86756574, 9293.16030173, 4902.96088326, 14725.62603104,
       14265.61948852, 9710.88151382, 34700.30707048, 6428.57784804,
       8599.07387595, 20697.64419021])

```

```

df1=pd.DataFrame({"Actual_value":y_test,"Predicted_value":y_pred,'Difference':y_test-y_pred})
df1

```

	Actual_value	Predicted_value	Difference
0	30760.000	25263.711557	5496.288443
1	17859.167	17510.601009	348.565991
2	9549.000	9440.299616	108.700384
3	11850.000	13532.184963	-1682.184963
4	28248.000	26109.694749	2138.305251
...
57	11845.000	9710.881514	2134.118486
58	37028.000	34700.307070	2327.692930
59	5389.000	6428.577848	-1039.577848
60	9233.000	8599.073876	633.926124
61	17199.000	20697.644190	-3498.644190

```
[62 rows x 3 columns]
```

Performance Evaluation

```

# 1. MAE
from sklearn.metrics import mean_absolute_error
print('MAE is ',mean_absolute_error(y_test,y_pred))

MAE is  2196.0531111239407

```

2. MAPE

```
from sklearn.metrics import mean_absolute_percentage_error  
print('MAE is ',mean_absolute_percentage_error(y_test,y_pred))
```

MAE is 0.19661972274847273

3. MSE

```
from sklearn.metrics import mean_squared_error  
print('MAE is ',mean_squared_error(y_test,y_pred))
```

MAE is 11091618.680797806

4. RMSE

```
from sklearn.metrics import mean_squared_error  
mse=mean_squared_error(y_test,y_pred)  
rmse=np.sqrt(mse)  
rmse
```

3330.4081853126963

5. R2_score

```
from sklearn.metrics import r2_score  
print(r2_score(y_test,y_pred))
```

0.8399116957802601